

# Optimisation for parameter estimation

(MATH10093: L03)

- ▶ In statistics, we often want to find the combination  $\hat{\boldsymbol{\theta}}$  of parameter values  $\boldsymbol{\theta} = \{\theta_1, \dots, \theta_m\}$  that maximises the *likelihood function*  $L(\mathbf{y}; \boldsymbol{\theta})$ .
- ▶ In special cases, we can use analysis to find closed form expressions for  $\hat{\boldsymbol{\theta}}$ .  
Example: If  $\mathbf{y} = \{y_1, \dots, y_n\}$  are independent observations of  $y_i \sim \mathcal{N}(\mu, \sigma^2)$ , the likelihood is


$$L(\mathbf{y}; \mu, \sigma^2) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left[ -\frac{(y_i - \mu)^2}{2\sigma^2} \right]$$

which is maximised by  $\hat{\mu} = \frac{1}{n} \sum_{i=1}^n y_i$  and  $\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{\mu})^2$ .


## Optimisation for parameter estimation

- ▶ For more complicated models, it may be difficult or impossible to find a solution by hand.
- ▶ Example: Change the example model by letting  $\sigma^2$  depend on a covariate, e.g. as  $\log(\sigma_i) = \theta_1 + x_i\theta_2$ .  
Now, the  $y_i$  are independent realisations from  $y_i \sim \text{N}(\mu, \sigma_i^2)$ .
- ▶ This *log-linear* model for the standard deviation doesn't provide a simple/analytical maximum likelihood solution to finding  $\hat{\theta}_1$  and  $\hat{\theta}_2$ .
- ▶ We need *numerical* optimisation methods!
- ▶ We usually convert the likelihood function into a related *target function*  $f(\theta)$  that is then *minimised*.


Target function  $f(\sigma) = L(\mathbf{y}; \sigma)$ ; has inflexion points




Target function  $f(\sigma) = \log L(\mathbf{y}; \sigma)$ ; is very skewed



Target function  $f(\theta) = \log L(\mathbf{y}; \sigma = e^\theta)$ ; have theory for *minimisation*



Target function  $f(\theta) = -\log L(\mathbf{y}; \sigma = e^\theta)$ ; the *negative log-likelihood*



# Searching for a minimum

Let  $\mathbf{g}(\boldsymbol{\theta})$  be the gradient vectors of  $f(\cdot)$ , i.e.  $g_j(\boldsymbol{\theta}) = \frac{\partial f(\boldsymbol{\theta})}{\partial \theta_j}$ . At a minimum,  $\|\mathbf{g}(\boldsymbol{\theta})\| = 0$ .

## Local minimum search algorithm

Start at some  $\boldsymbol{\theta}^{[0]}$ , and iterate over  $\boldsymbol{\theta}^{[k]}$ ,  $k = 0, 1, 2, \dots$ :

1. Find a *descent direction* vector  $\mathbf{d}^{[k]}$  from  $\boldsymbol{\theta}^{[k]}$ . This means that  $\mathbf{g}(\boldsymbol{\theta}^{[k]})^\top \mathbf{d}^{[k]} < 0$ .
2. Perform a *line search*, by finding a *step scaling*  $\alpha_k > 0$  such that the new  $f$  value is sufficiently improved:  
$$f(\boldsymbol{\theta}^{[k]} + \alpha_k \mathbf{d}^{[k]}) < f(\boldsymbol{\theta}^{[k]}) + \epsilon \alpha_k \mathbf{g}(\boldsymbol{\theta}^{[k]})^\top \mathbf{d}^{[k]}$$
for some fixed  $0 < \epsilon < 1$  (can be small, e.g.  $10^{-3}$ )
3. Let  $\boldsymbol{\theta}^{[k+1]} = \boldsymbol{\theta}^{[k]} + \alpha_k \mathbf{d}^{[k]}$

Terminate the iteration when either no improvement is found or we have reached a minimum:

$k > \text{maximum allowed iteration steps,}$

$$\|\boldsymbol{\theta}^{[k+1]} - \boldsymbol{\theta}^{[k]}\| < \text{tol}_x,$$

$$f(\boldsymbol{\theta}^{[k]}) - f(\boldsymbol{\theta}^{[k+1]}) < \text{tol}_f, \text{ or}$$

$$\|\mathbf{g}(\boldsymbol{\theta}^{[k]})\| < \text{tol}_g.$$

## A simple and practical line search method

Given a valid descent direction  $\mathbf{d}^{[k]}$  we know that an  $\alpha_k > 0$  exists, such that the function value at  $\boldsymbol{\theta}^{[k]} + \alpha_k \mathbf{d}^{[k]}$  is lower than at the starting point,  $f(\boldsymbol{\theta}^k)$ .

A simple *inexact line search* method:

1. Let  $\alpha_k = 1$ .
2. Stop if  $f(\boldsymbol{\theta}^{[k]} + \alpha_k \mathbf{d}^{[k]}) < f(\boldsymbol{\theta}^{[k]}) + \epsilon \alpha_k \mathbf{g}(\boldsymbol{\theta}^{[k]})^\top \mathbf{d}^{[k]}$ .
3. Otherwise, divide  $\alpha_k$  by 2, and go back to step 2.

Provided that  $0 < \epsilon < 1$ , this iteration will terminate.



## Gradient descent direction with adaptive step length

The most basic descent direction is the reverse gradient:

$$\mathbf{d}^{[k]} = -\gamma_k \mathbf{g}(\boldsymbol{\theta}^{[k]}) / \|\mathbf{g}(\boldsymbol{\theta}^{[k]})\|,$$

where  $\gamma_k > 0$  is the *proposed step length*. A fixed  $\gamma_k$  is inefficient.

In the first step, let  $\gamma_0 = 1$ , and for  $k = 1, 2, \dots$ ,

if  $\alpha_{k-1} = 1$  (the proposed length was OK), let  $\gamma_k = 3\gamma_{k-1}/2$  (try a longer step next time),  
otherwise  $\gamma_k = \alpha_{k-1}\gamma_{k-1}$  (reuse the latest actual accepted step length).

## Newton optimisation

A second order Taylor series approximation of  $f(\cdot)$  contains useful information about the size shape of the target function.

Let  $\mathbf{H}(\boldsymbol{\theta})$  be the second order derivative matrix (or *Hessian*) of  $f(\cdot)$ , with  $H_{ij}(\boldsymbol{\theta}) = \frac{\partial^2 f(\boldsymbol{\theta})}{\partial \theta_i \partial \theta_j}$ .

### Newton search direction

The quadratic approximation of  $f(\cdot)$ ,

$$f(\boldsymbol{\theta} + \mathbf{d}) \approx f(\boldsymbol{\theta}) + \mathbf{g}(\boldsymbol{\theta})^\top \mathbf{d} + \frac{1}{2} \mathbf{d}^\top \mathbf{H}(\boldsymbol{\theta}) \mathbf{d},$$

constructed at  $\boldsymbol{\theta}^{[k]}$ , is minimised by taking a step

$$\mathbf{d}^{[k]} = -\mathbf{H}(\boldsymbol{\theta}^{[k]})^{-1} \mathbf{g}(\boldsymbol{\theta}^{[k]})$$

if the Hessian is positive definite (has strictly positive eigenvalues).

Bad news: The Hessian  $\mathbf{H}$  is not always positive definite far away from the minimum.

Good news: Replacing  $\mathbf{H}$  by *any* positive definite matrix leads to a descent direction.

Solution: Find practical approximations to  $\mathbf{H}$  that are positive definite. Example: *BFGS*

## Quasi-Newton method example: BFGS

One of the most popular *Quasi-Newton* methods requires only function values and gradients.

### Broyden-Fletcher-Goldfarb-Shanno (BFGS)

The BFGS method can be formulated to work directly with an approximation to  $\mathbf{H}(\boldsymbol{\theta})^{-1}$ , removing the need to solve a linear system to find the descent direction.

1. Let  $\mathbf{B}^{[0]}$  be a guess of  $\mathbf{H}(\boldsymbol{\theta}^{[0]})^{-1}$ .
2. For each  $k$ , compute the step  $\mathbf{a}_k = \alpha_k \mathbf{d}^{[k]}$  from line search using the search direction  $\mathbf{d}^{[k]} = -\mathbf{B}^{[k]} \mathbf{g}(\boldsymbol{\theta}^{[k]})$ .
3. Compute  $\mathbf{b}_k = \mathbf{g}(\boldsymbol{\theta}^{[k]} + \mathbf{a}_k) - \mathbf{g}(\boldsymbol{\theta}^{[k]})$ , and update the  $\mathbf{B}$  matrix:

$$\mathbf{B}^{[k+1]} = \mathbf{B}^{[k]} + \frac{\mathbf{a}_k^\top \mathbf{b}_k + \mathbf{b}_k^\top \mathbf{B}^{[k]} \mathbf{b}_k}{(\mathbf{a}_k^\top \mathbf{b}_k)^2} \mathbf{a}_k \mathbf{a}_k^\top - \frac{\mathbf{B}^{[k]} \mathbf{b}_k \mathbf{a}_k^\top + \mathbf{a}_k \mathbf{b}_k^\top \mathbf{B}^{[k]}}{\mathbf{a}_k^\top \mathbf{b}_k}$$

The equations guarantee that  $\mathbf{B}^{[k]}$  stays positive definite.

The initial  $\mathbf{B}^{[0]}$  is often chosen to be either proportional to an identity matrix, or a diagonal matrix based on the diagonal elements of  $\mathbf{H}(\boldsymbol{\theta}^{[0]})$ , which costs only around twice as much as a single gradient calculation.

## Computational cost considerations

- ▶ Computing  $f$ ,  $g$ , and  $H$  is often expensive.
- ▶ When using finite differences, the cost of  $g$  is proportional to  $m$ , and the cost of  $H$  is proportional to  $m^2$ .
- ▶ We want to balance the cost per iteration with the number of iterations required to reach the minimum.
- ▶ It's usually not worth computing the actual second order derivatives, unless we can find a closed form positive definite Hessian approximation.
- ▶ From the theory of negative log-likelihood functions it is known that the *expected Hessian*,  $\widetilde{H}(\theta) = E_{y|\theta}[H(\theta)]$  is always positive definite. Using this in place of the *observed Hessian*  $H(\theta)$  is called *Fisher Scoring*.
- ▶ Conclusion:
  - ▶ For smooth target functions with cheap gradients, BFGS or Fisher Scoring is preferable
  - ▶ For less smooth target functions or expensive gradients, the Simplex method is preferable (robust and uses only  $f$  values, see Computer Lab 2; this is also the default method in `optim()` in R).
- ▶ In Computer Lab 2, you will experiment with different target functions and optimisation methods in a graphical interactive R tool.

## Fisher Scoring example

Let  $y_i \sim N(\mu(\boldsymbol{\theta}), \sigma(\boldsymbol{\theta})^2)$  (independent),  $\mu(\boldsymbol{\theta}) = \theta_1$ ,  $\sigma(\boldsymbol{\theta}) = e^{\theta_2}$ . The negative log-likelihood is

$$f(\theta_1, \theta_2) = \frac{n}{2} \log(2\pi) + n\theta_2 + \frac{1}{2e^{2\theta_2}} \sum_{i=1}^n (y_i - \theta_1)^2$$

The gradient elements are

$$g_1(\theta_1, \theta_2) = -\frac{1}{e^{2\theta_2}} \sum_{i=1}^n (y_i - \theta_1), \quad g_2(\theta_1, \theta_2) = n - \frac{1}{e^{2\theta_2}} \sum_{i=1}^n (y_i - \theta_1)^2$$

The observed and expected Hessian elements are

$$\begin{aligned} H_{11}(\theta_1, \theta_2) &= \frac{n}{e^{2\theta_2}}, & H_{12}(\theta_1, \theta_2) &= \frac{2}{e^{2\theta_2}} \sum_{i=1}^n (y_i - \theta_1), & H_{22}(\theta_1, \theta_2) &= \frac{2}{e^{2\theta_2}} \sum_{i=1}^n (y_i - \theta_1)^2, \\ \tilde{H}_{11}(\theta_1, \theta_2) &= \frac{n}{e^{2\theta_2}}, & \tilde{H}_{12}(\theta_1, \theta_2) &= 0, & \tilde{H}_{22}(\theta_1, \theta_2) &= 2n. \end{aligned}$$

The expected Hessian is diagonal with positive elements (so clearly positive definite) and much cheaper to compute, since the observations are not involved!